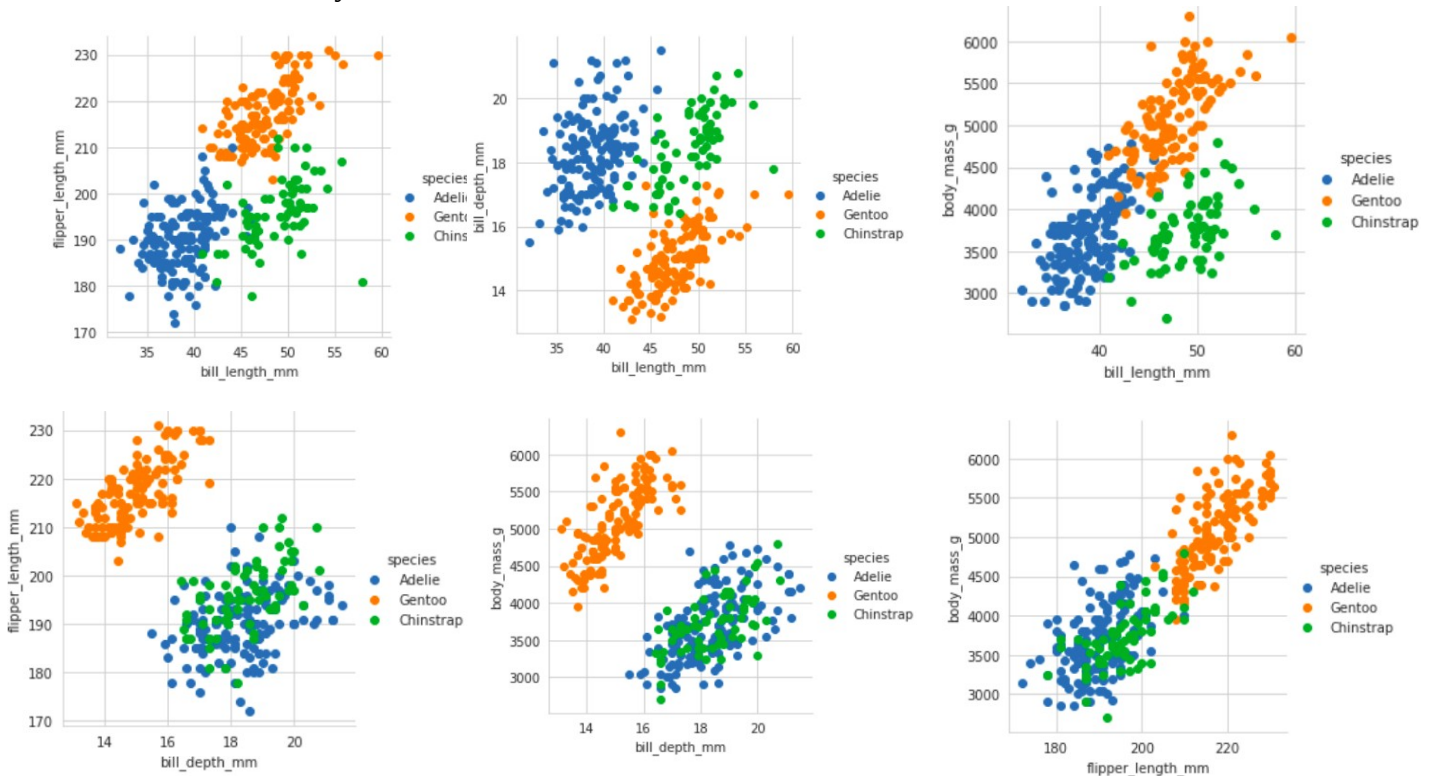


PRML Lab-2
Abu Shahid (B20CS003)- [Link to the File](#)

Question 1
Subtask 1

- Values of few features were missing and therefore, corresponding rows were dropped.
- Our data was 7-dimensional and therefore visualization was not straightforward.
- Just to get an insight of the classes, various features were plotted against one another and visualized one by one.



- As seen by the previous plot, even by using two features, the classes are quite separable for few select feature pairs and we can expect to have a well defined decision boundary. (Such an exercise overkill here, would have been helpful for feature selection to train our model had the dataset been enormous)
- On inspection of the dataset, **year**, **island** and **sex** were found to be categorical labels and one-hot encoded directly(skipping the intermediate step of converting them to categorical codes). Sex was also one-hot-encoded. The species being the target variable was left untouched.
- The dataset was then split into 80:20 ratio of train and test set.

Subtask-2

- Implementing gini index calculation as per: $\text{gini index} = 1 - \sum (p_i^2)$. (the function for calculation of information gain was not coded explicitly)

```
def gini_index(data):
    classes= np.unique(data)
    gini=1

    for c_i in classes:
        p= len( data[data==c_i])/len(data)
        gini-=p**2
    return gini
```

Subtask-3

- **con_to_cat()** function was declared and the pseudocode of the same is:

> def con_to_cat(data):

> split={} //dictionary -to store the left and right data after the split is made
-to store the feature which was analysed for the split
-the value of the feature(threshold) used for splitting
-the resulting gini index for the split

> a nested iterater to loop through all the features and all the values of the features, calculating all the possible splits and returning the case with best split.

```
[320] def con_to_cat(data):
    sample_count=data.shape[0]
    feature_count=data.shape[1]

    max_gini=-999999999999
    split={}

    for feature in range(feature_count):
        values=data[:, feature]
        values=np.unique(values)
        for threshold in values:
            data_left=np.array( [row for row in data if row[feature]<= threshold])
            data_right=np.array( [row for row in data if row[feature]> threshold])

            if( len(data_left)>0 and len(data_right)>0):
                y, left_y, right_y= data[:, -1], data_left[:, -1], data_right[:, -1]

                wl=len(left_y)/len(y)
                wr=len(right_y)/len(y)
                gini= gini_index(y)- wl*gini_index(left_y) - wr*gini_index(right_y)
                if gini> max_gini and feature!=12: #####second condition added idky
                    max_gini= gini
                    split['gini']=max_gini
                    split['threshold']= threshold
                    split['left']= data_left
                    split['right']=data_right
                    split['feature']=feature
                    #print(feature)

    return split
```

Subtask 4 and 5

- Our **con_to_cat()** function itself gets the attribute for the best split and then make the split.
- In our implementation, **con_to_cat()** is continuously called recursively called and splits are made till
 - data to be split is > min_samples_leaf
 - depth < max_depth
 - and information gain for the split> 0
- This was majorly carried out by the **tree()** function which gets automatically called on calling **fit()**.
- The **tree()** implements the **con_to_cat()** and **gini()** function and makes the decision tree as per above conditions.
- Referenced from: https://github.com/Suji04/ML_from_Scratch/blob/master/decision%20tree%20classification.ipynb


Question 2

Subtask 1

- The data was analysed and it was a complete dataset with no missing values.
- Next we calculated the correlation of different features with our target variable Y1.
 - X5 0.889431, X1 0.622272, X3 0.455671, X7 0.269841, X8 0.087368, X6 -0.002587, X2 -0.658120, X4 -0.861828
- Such an exercise, here however unnecessary would help us in feature selection had the dataset been gigantic.
- Dataset source was inspected and it was found that **X6** was a categorical label and therefore was one-hot-encoded. Data was then modified with the one-hot labels appended and original X6 was dropped.
- Finally data was split into train and test set, first in the ration 80:20
- Then it was further split into train and validation set, in the ration 87.5 : 12.5

Subtask 2

- To get an intuition of how our model behaves on changing hyperparameters; we vary one hyperparameter at a time, training on the training set and calculating MSE on the validation set. The results are plotted.



- On a rough note MSE increases with increase in **min_sample_split** and **min_sample_leaves** and decreases with increases in **max_depth**.
- This is so because, as the depth inceases, our model can made diverse threshold conditions to fit more of the data and define complex decision boundaries.
- And as the min_sample_split decreases, our model has more freedom to reduce the entropy of a leaf and increase the information gain by further splitting.
- To find the best combination of parameters, a function is defined with trains our model over varying hyperparameters and returns the combination which gives the best MSE.

```
def best_params(X_train, y_train, X_val, y_val, best_mse):  
    for depth in range(2,30):  
        for sample_split in range(2,12):  
            for sample_leaf in range(1,10):
```

```
    {'depth': 8, 'sample_split': 3, 'sample_leaf': 1}
```

Subtask 3


- Using the parameters we got from the previous exercise, we train a DecisionTreeRegressor using the same parameters and to validate its performance, we perform k-fold cross validation. (k=5). The results are as:


```
array([0.99693291, 0.99685029, 0.99705815, 0.99614764, 0.99543896])
```

- Not only are the validation scores consistent, ensuring that our model is not over-fitting but also they are high: ensuring that our model has indeed learned.
- Finally we calculate the mean squared error between the predicted and the ground-truth values in the test data for our best model.

```
0.3506566629935371
```

- At last, tree was plotted:



and some decision nodes for the data that is central in our dataset;

